

Course. Introduction to Machine Learning

Lecture 7. Lazy learning

Dr. Maria Salamó Llorente
Dept. Mathematics and Informatics,
Faculty of Mathematics and Informatics
University of Barcelona (UB)

Introduction to Machine Learning

Unsupervised Learning

Cluster Analysis

Factor Analysis

Visualization

K-Means,
Fuzzy C-means,
EM

PCA, ICA

Self Organized Maps (SOM) ,
Multi-Dimensional Scaling

Lazy Learning
(K-NN, IBL,
CBR)

Overfitting,
model selection and
feature selection

Kernel Learning

Ensemble Learning
(Trees,
Adaboost)

Perceptron,
SVM

Bias/Variance,
VC dimension,
Practical advice of how
to use learning
algorithms

Supervised Learning

Non Linear Decision

Linear Decision

Decision Learning Theory



1. Introduction to lazy learning
2. Nearest neighbor
3. Instance-based learning
4. K-Nearest neighbor
5. Case-based reasoning
6. Summary
7. References



UNIVERSITAT DE BARCELONA



Introduction to Lazy Learning

- Lazy learning delays generalization until a query is made to the system
- Advantages:
 - The target function is approximated locally
 - It can simultaneously solve multiple problems
 - It deals successfully with changes in the problem domain
 - Suitable for complex and incomplete problem domains
- Disadvantages:
 - It requires a large space to store the entire training
 - It may be slow for solving a problem but it has a fast training

In this lesson we will learn about:

- Nearest Neighbour
- Instance-based learning
- Case-based reasoning

{case, exemplar, instance, memory}

-based-

{learning, reasoning}



UNIVERSITAT DE BARCELONA



Nearest Neighbor

(NN)

Basic idea:

1. Get some example set of cases with known outputs
 - E.g. Diagnoses of infectious diseases by experts
2. When you see a new case, assign its output to be the same as the most similar known case
 - Your symptoms most resemble Mr. X
 - Mr. X had a flu
 - Ergo you have the flu

- There is a set of possible examples, $X = \{x_i\}$
- Each example is an n-tuple of attribute values, $\vec{x}_1 = \langle a_1, \dots, a_k \rangle$
- There is a target function that maps X onto set Y,
 $f: X \rightarrow Y$
- The data is a set of tuples <example, target function values>

$$D = \{\langle \vec{x}_1, f(\vec{x}_1) \rangle, \dots, \langle \vec{x}_m, f(\vec{x}_m) \rangle\}$$

- Find a hypothesis h such that

$$\forall \vec{x}, h(\vec{x}) \approx f(\vec{x})$$

- Task: Given some set of training data

$$D = \{\langle \vec{x}_1, f(\vec{x}_1) \rangle, \dots, \langle \vec{x}_m, f(\vec{x}_m) \rangle\}$$

and a query point \vec{x}_q , predict $f(\vec{x}_q)$

1. Find the nearest member of the data set to the query

$$\vec{x}_{nn} = \arg \min_{x \in D} (d(\vec{x}, \vec{x}_q))$$

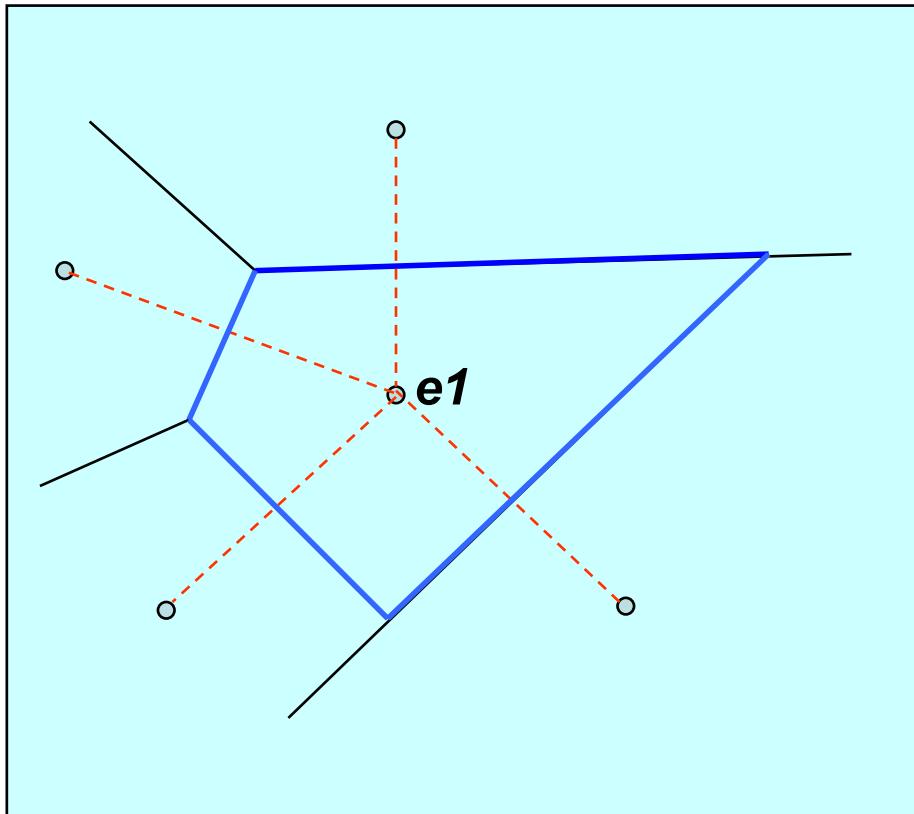
2. Assign the nearest neighbour output to the query

$$h(\vec{x}_q) = f(\vec{x}_{nn})$$

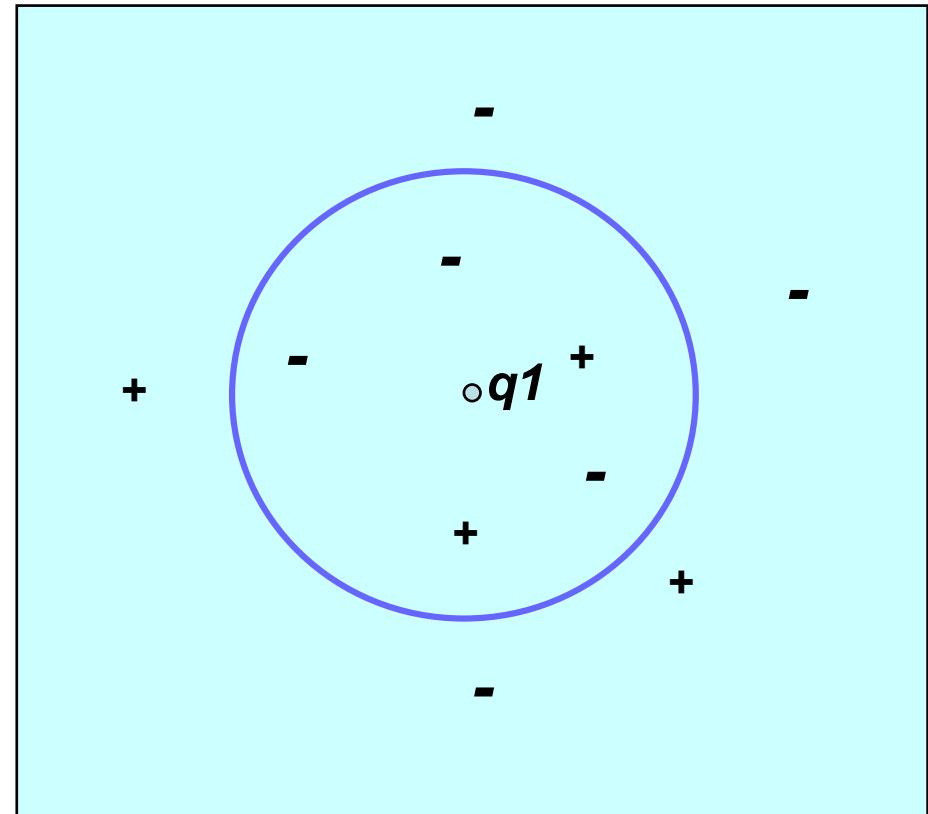
Our hypothesis

distance function

Voronoi diagram

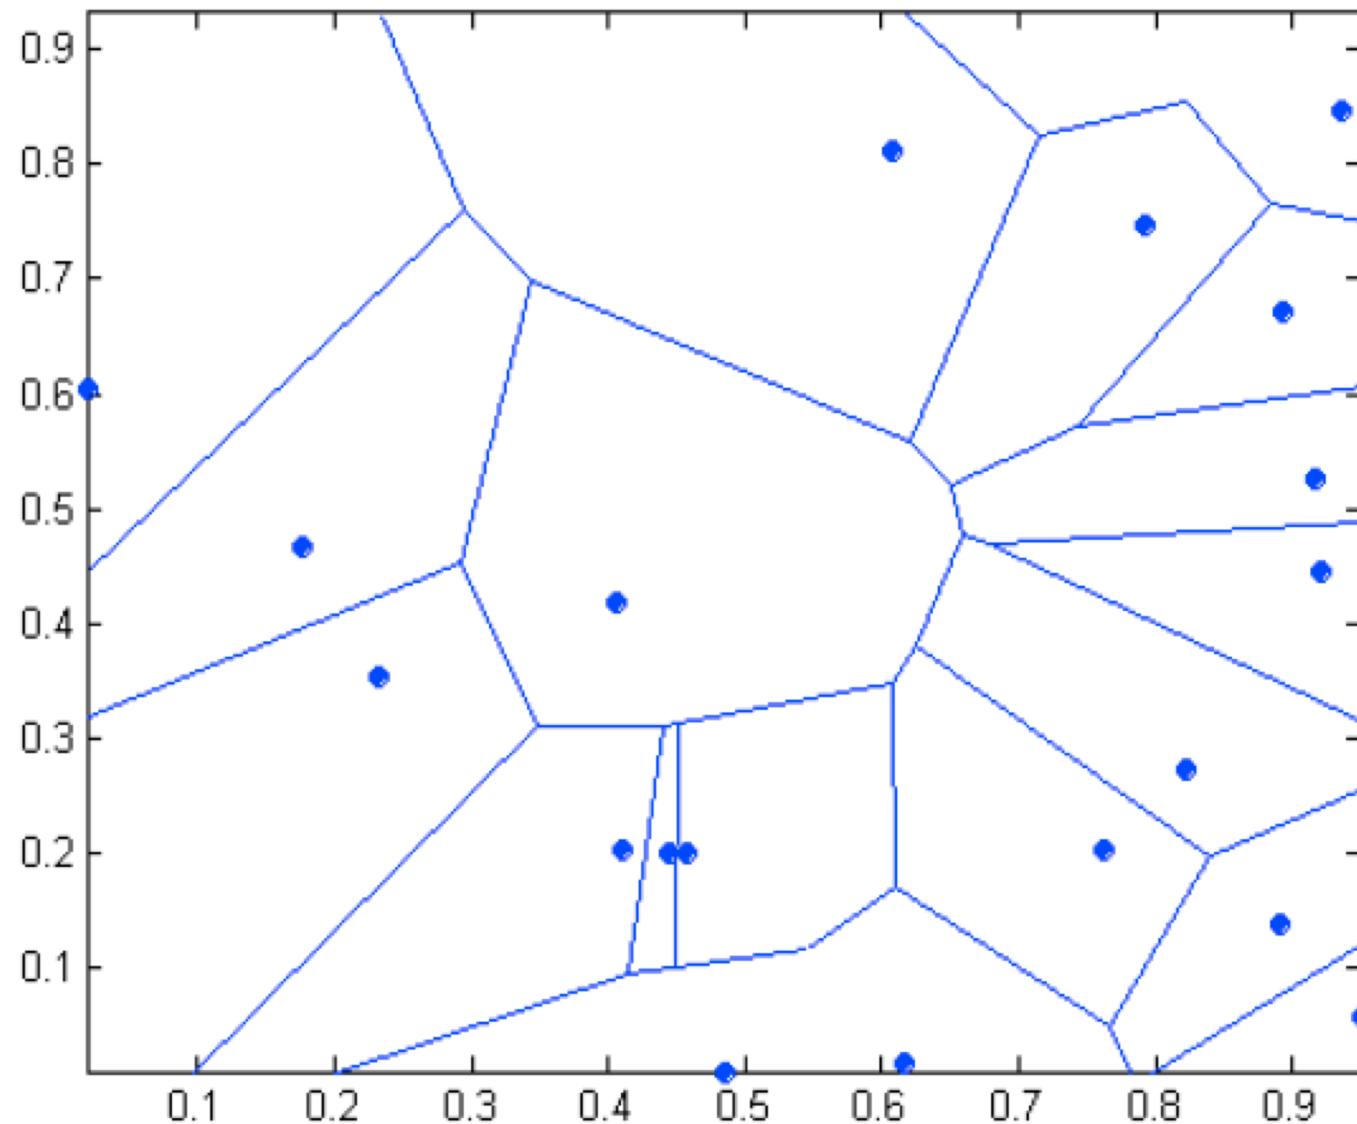


*1-nearest neighbor:
the concept represented by $e1$*



*5-nearest neighbors:
 $q1$ is classified as negative*

Voronoi diagram



Instance-based Learning

(IBL)

- A lazy learning algorithm
- One way of solving tasks of approximating **discrete or real valued** target functions
- It produces local approximations to the target function
- Key idea:
 - just store the training examples, D
 - when a test example is given then **find the closest** matches

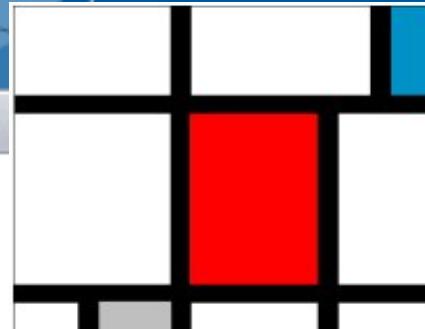
- A distance measure
 - Nearest neighbour: typically Euclidean
- Number of neighbours to consider
 - Nearest neighbour: one
- A weighting function (optional)
 - Nearest neighbour: unused (equal weights)
- How to fit with neighbours
 - Nearest neighbour: same output as nearest neighbour

K-Nearest Neighbor (KNN)

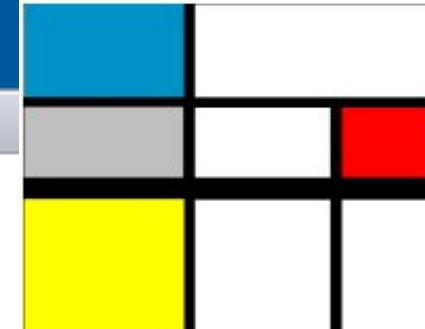
- A distance measure
 - typically Euclidean
- Number of neighbours to consider
 - k
- A weighting function (optional)
 - Unused (equal weights)
- How to fit with neighbours
 - Vote using K nearest neighbours (or take average, for regression)

- Instances map to points in \mathbb{R}^n
- Not more than say 20 attributes per instance
- Lots of training data
- Advantages:
 - Training is very fast
 - Can learn complex target functions
 - Don't lose information
- Disadvantages:
 - ??? (We will see them shortly...)

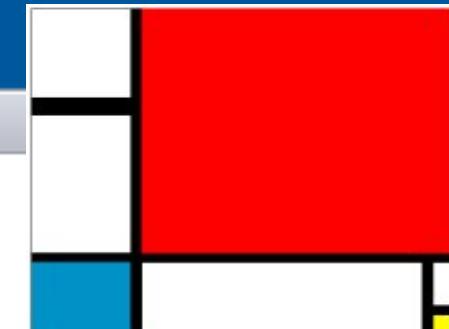
Mondrian Example



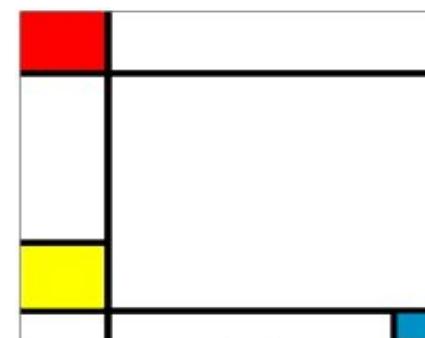
one



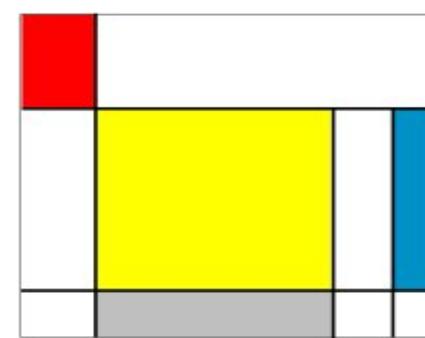
two



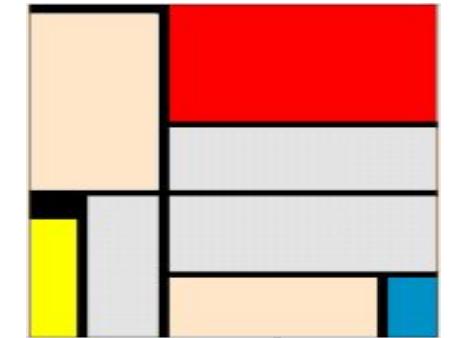
three



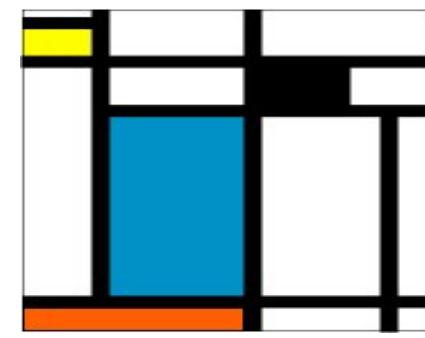
four



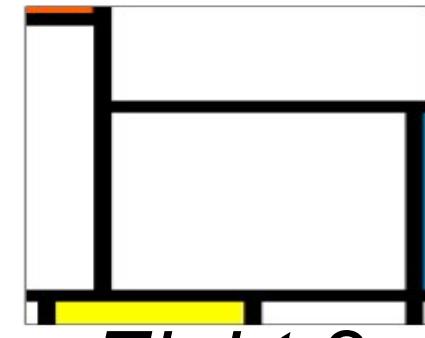
five



six



seven



Eight ?

Training data

Number	Lines	Line types	Rectangles	Colours	Mondrian?
1	6	1	10	4	No
2	4	2	8	5	No
3	5	2	7	4	Yes
4	5	1	8	4	Yes
5	5	1	10	5	No
6	6	1	8	6	Yes
7	7	1	14	5	No

Test instance x_q

Number	Lines	Line types	Rectangles	Colours	Mondrian?
8	7	2	9	4	

Keep data in normalised form

One way to normalise the data $a_r(x)$ to $a'_r(x)$ is

$$x_t' \equiv \frac{x_t - \bar{x}_t}{\sigma_t}$$

\bar{x}_r *≡ mean of t^{th} attributes*

σ_t *≡ standard deviation of t^{th} attributes*

Keep data in normalised form

Another way to normalise the data is rescaling the range of features to the range [0, 1] or [-1 , 1]

$$x_t' \equiv \frac{x_t - \min}{\max - \min}$$

$\min \equiv \min$ of t^{th} attributes

$\max \equiv \max$ of t^{th} attributes

Normalised training data

Number	Lines	Line types	Rectangles	Colours	Mondrian?
1	0.632	-0.632	0.327	-1.021	No
2	-1.581	1.581	-0.588	0.408	No
3	-0.474	1.581	-1.046	-1.021	Yes
4	-0.474	-0.632	-0.588	-1.021	Yes
5	-0.474	-0.632	0.327	0.408	No
6	0.632	-0.632	-0.588	1.837	Yes
7	1.739	-0.632	2.157	0.408	No

Test instance (x_q)

Number	Lines	Line types	Rectangles	Colours	Mondrian?
8	1.739	1.581	-0.131	-1.021	

Example	Distance of test instance to the example	Mondrian?
1	2.517	No
2	3.644	No
3	2.395	Yes
4	3.164	Yes
5	3.472	No
6	3.808	Yes
7	3.490	No

Classification

1-NN Yes

3-NN Yes

5-NN No

7-NN No

- We might want to weight nearer neighbours more heavily

$$f(\mathbf{x}_q) := \frac{\sum_{i=1}^k w_i f(\mathbf{x}_i)}{\sum_{i=1}^k w_i} \text{ where } w_i = \frac{1}{d(\mathbf{x}_q, \mathbf{x}_i)^2}$$

- Then it makes sense to use *all training examples* instead of just k (Stepard's method)

- Advantages
 - Fast training (**it's a lazy algorithm**)
 - Learn complex functions easily
 - Do not lose information
- Disadvantages
 - Slow at query time
 - Needs lots of storage
 - Easily fooled by irrelevant attributes
 - Use statistical methods to remove irrelevant ones (e.g., PCA or cross-validation)

- Nearest neighbour easily misled when X high-dim
- Low-dimensional intuitions do not extend to high dim
- **The curse of dimensionality**
 - When dimensionality increases, the volume of the space increases so fast that the available data becomes sparse.
 - The number of features is too large relative to the number of training samples

This makes a poor generalization ability of the classifier.

Case-Based Reasoning

(CBR)

- *In this course we focus on the basic principle rather than on specific applications or tools*
- Briefly
 - CBR is an advanced instance-based learning applied to more complex instance objects
 - Objects may include complex structural descriptions of cases and adaptation rules
 - The power comes from the organisation and content of the cases themselves

Faced this situation before?

- Oops the car stopped.
 - What could have gone wrong?
- Aah.. Last time it happened, there was no petrol.
 - Is there petrol?
 - Yes.
 - Oh but wait I remember the tyre was punctured (ban bocor)

This is the normal thought process of a human when faced with a problem which is similar to a problem he/she had faced before.

How do we solve problems?

- By knowing the steps to apply
 - from symptoms to a plausible diagnosis
- But not always applying causal knowledge
 - diseases cause symptoms
 - symptoms do not cause diseases!
- How does an expert solve problems?
 - uses same “book learning” as a novice
 - but quickly selects the right knowledge to apply
- Heuristic knowledge (“rules of thumb”)
 - “I don’t know why this works but it does and so I’ll use it again!”
 - difficult to elicit

- By remembering how we solved a similar problem in the past
- This is Case Based Reasoning (CBR)!
 - memory-based problem-solving
 - re-using past experiences
- Experts often find it easier to relate stories about past cases than to formulate rules

- **Medicine**
 - doctor remembers previous patients especially for rare combinations of symptoms
- **Law**
 - English/US law depends on precedence
 - case histories are consulted
- **Management**
 - decisions are often based on past rulings
- **Financial**
 - performance is predicted by past results

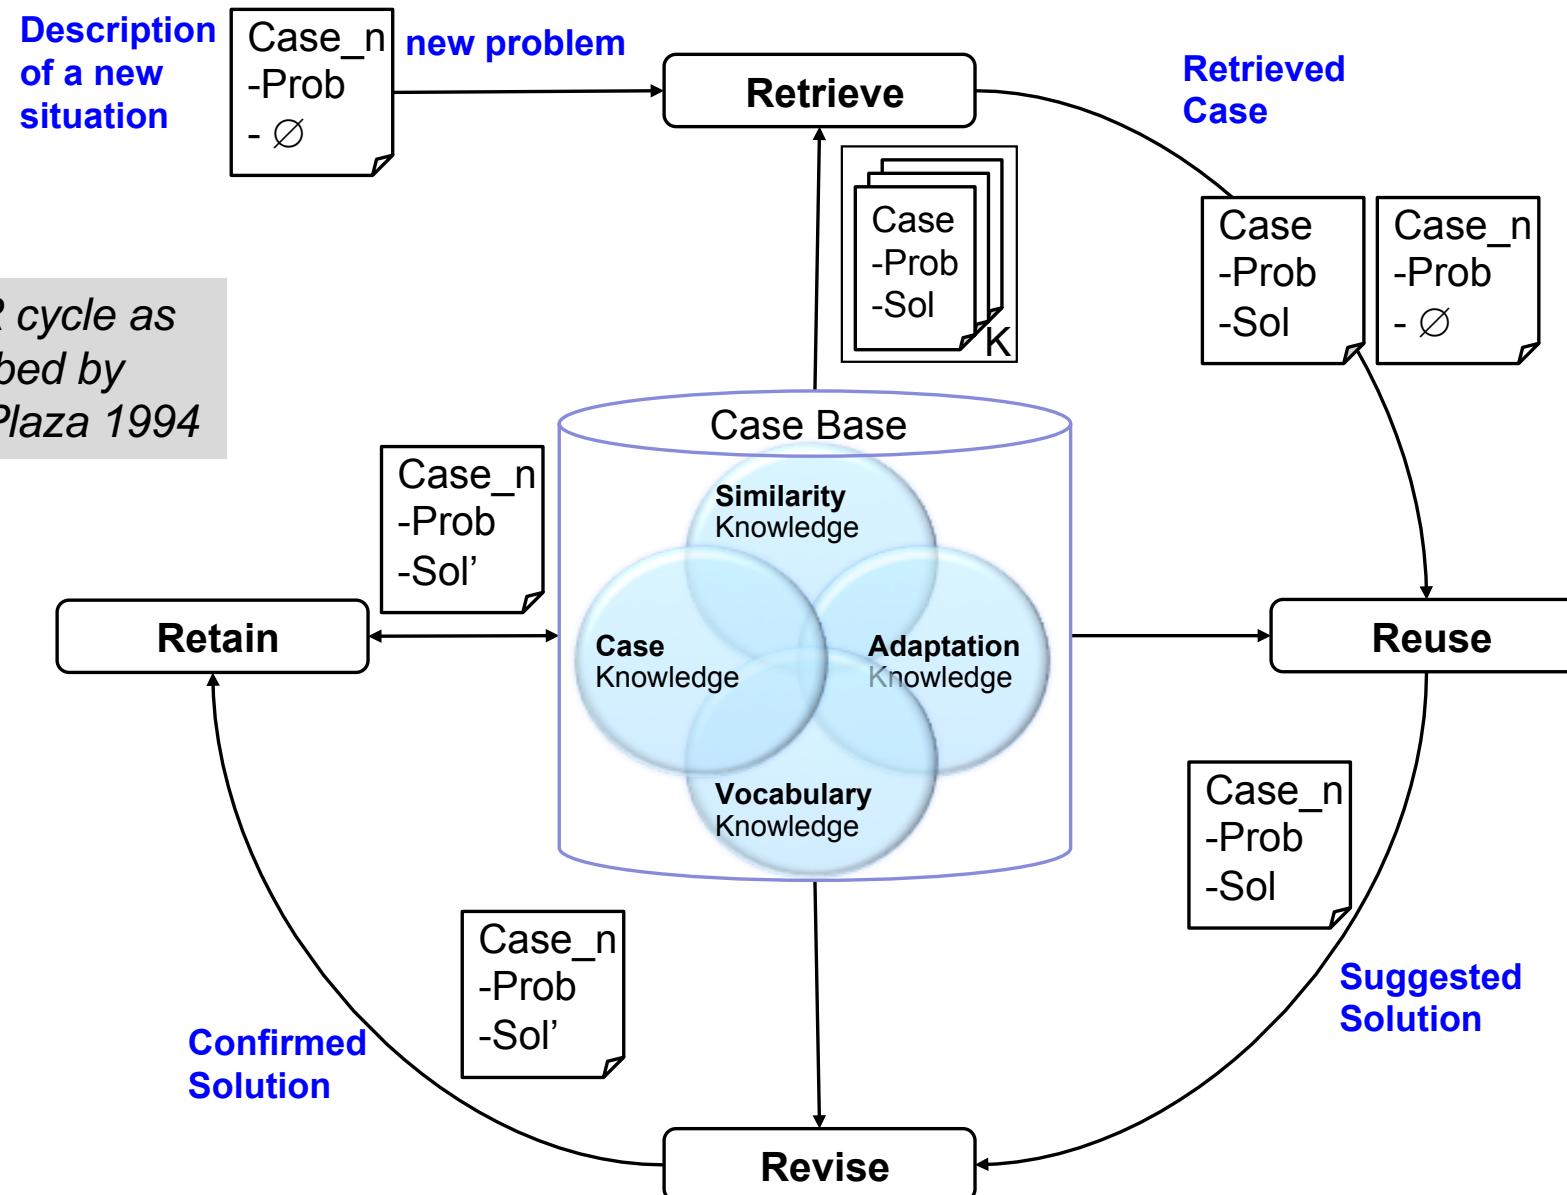
What is Case-based Reasoning

- Case-based reasoning is [...] **reasoning by remembering** – [Leake, 1996]
- A case-based reasoner solves new problems by adapting solutions that were used to solve old problems - [Riesbeck & Schank, 1989]
- Case-based reasoning is both [...] the ways people use cases to solve problems and the ways we can make machines use them – [Kolodner, 1993]
- Case-based reasoning is a recent approach to problem solving and learning [...] - [Aamodt & Plaza, 1994]

- A methodology to model human reasoning and thinking
- A methodology for building intelligent computer systems
- CBR in a nutshell:
 - Store previous experience (cases) in memory
 - To solve new problems:
 - Retrieve similar experience about similar situations from the memory
 - Reuse the experience in the context of the new situation : complete or partial reuse, or adapt according to differences
 - Store new experience in memory (**learning**)

- Roots of CBR is found in the works of Roger Shank on dynamic memory in 1983
- Other trails into the CBR field has come from
 - Analogical reasoning in 1990
 - Problem solving and experimental learning within philosophy and psychology
- The first CBR system, CYRUS developed by Janet Kolodner at Yale university (1983)

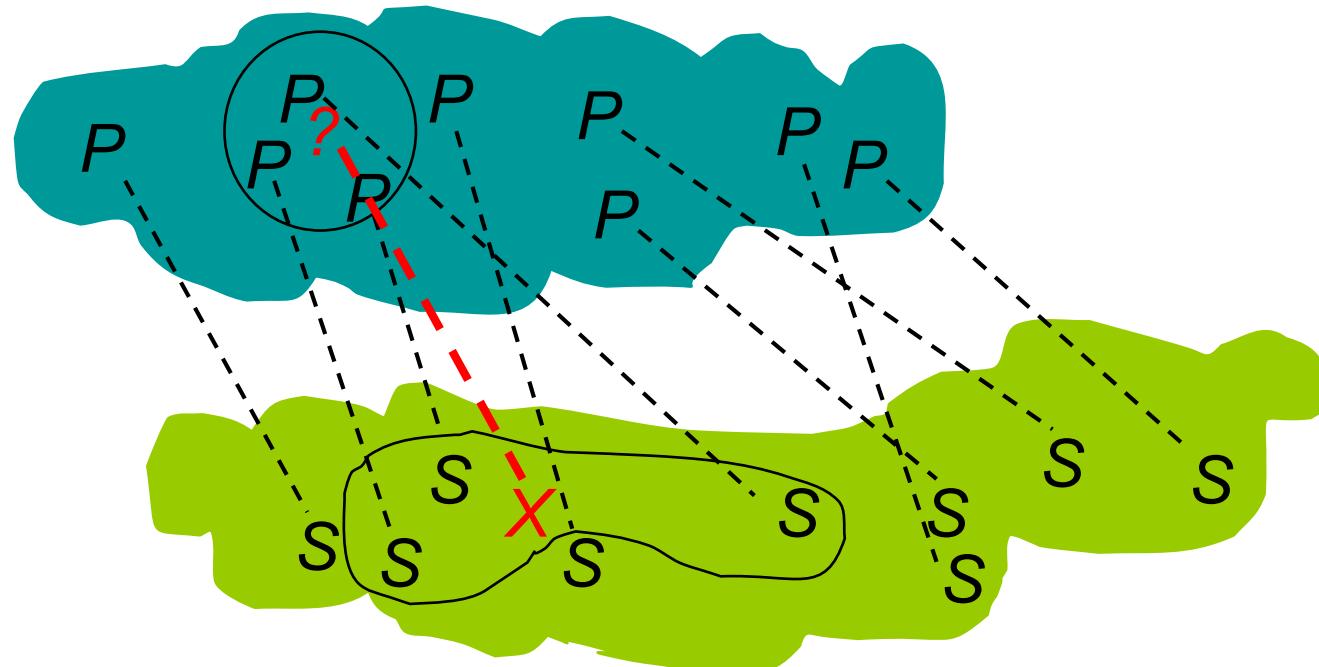
CBR cycle



- Case-base
 - database of previous cases (experience)
- Retrieval of relevant cases
 - index for cases in library
 - matching most similar case(s)
 - retrieving the solution(s) from these case(s)
- Adaptation of solution
 - alter the retrieved solution(s) to reflect differences between new case and retrieved case(s)

Main CBR Assumption

- New problem can be solved by
 - retrieving similar problems
 - adapting retrieved solutions
- Similar problems have similar solutions

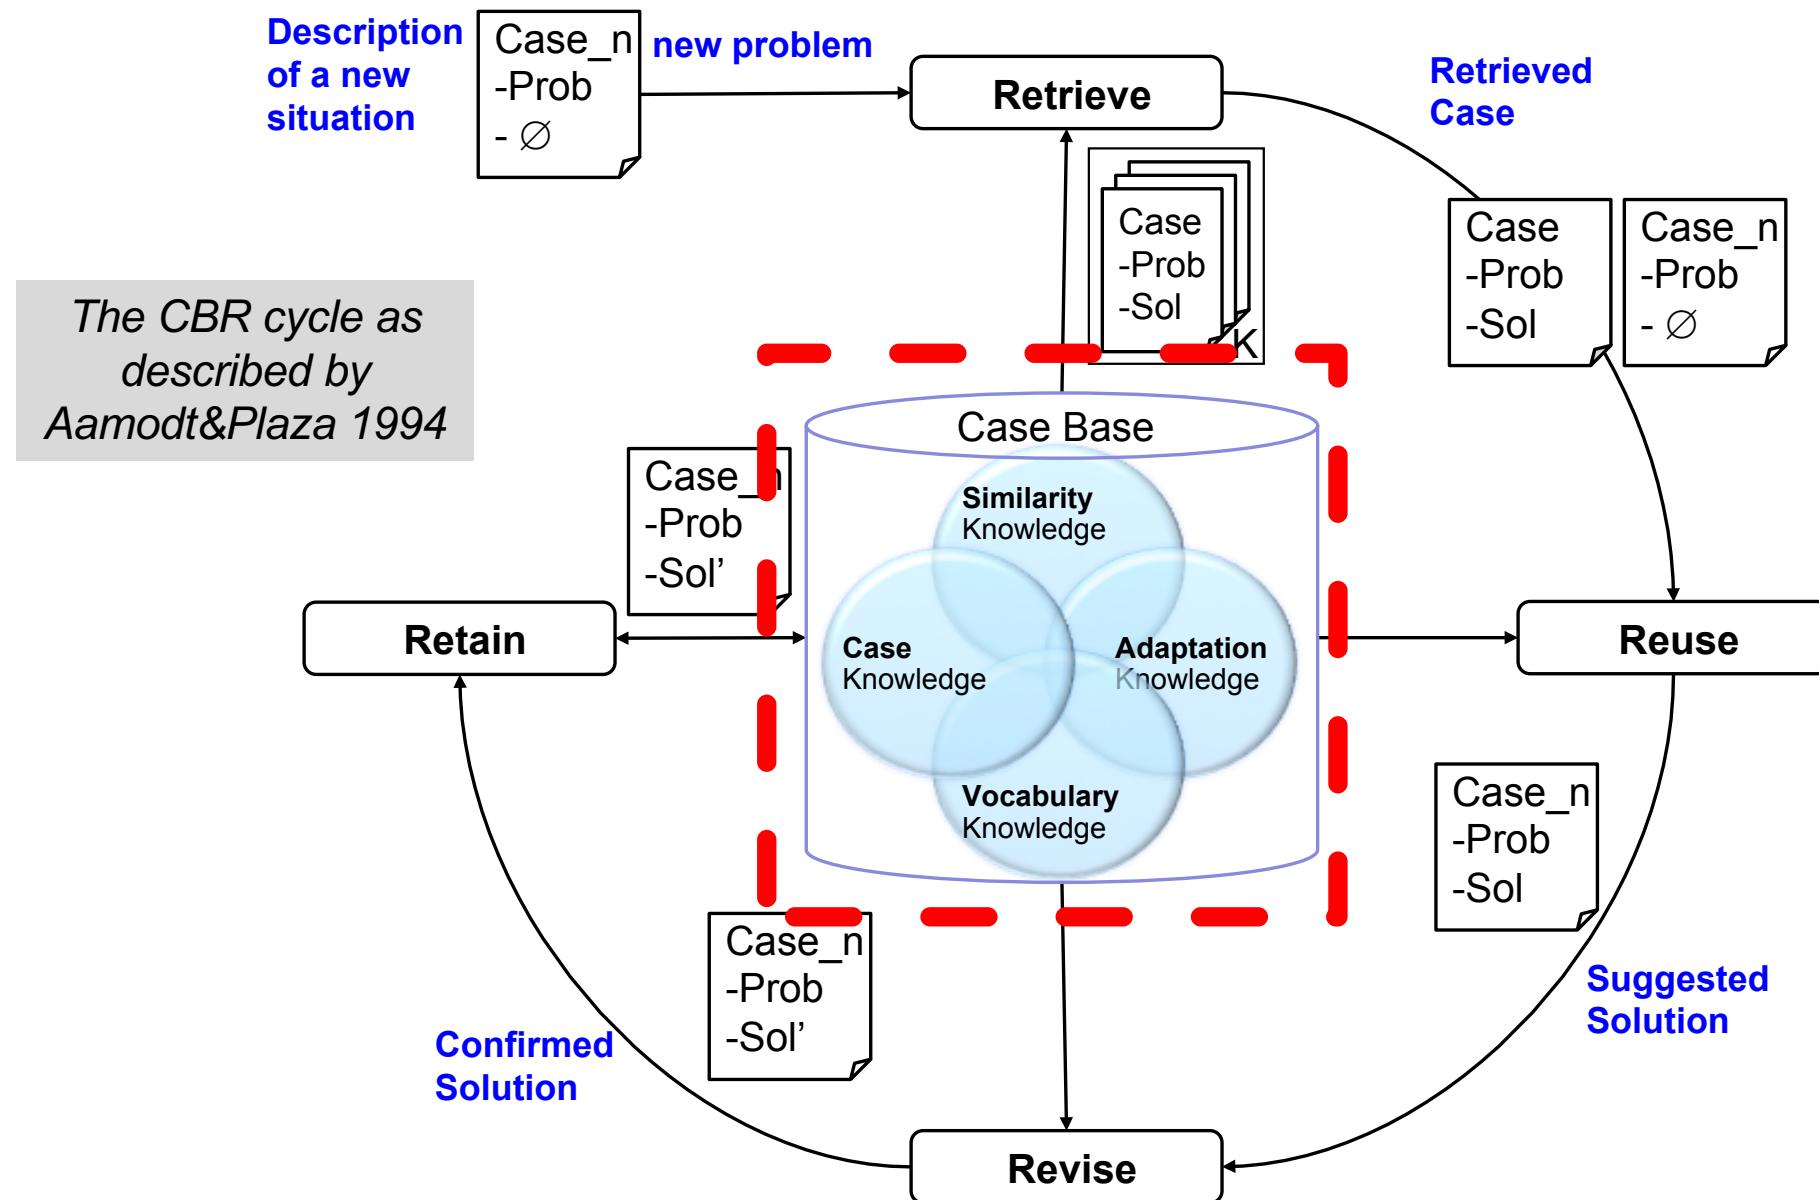


- The **main assumption** is that:
 - *Similar problems have similar solutions:*
 - e.g., an aspirin can be taken for any mild pain
- More assumptions:
 - **The world is a regular place:** what holds true today will probably hold true tomorrow
 - (e.g., if you have a headache, you take aspirin, because it has always helped)
 - **Situations repeat:** if they do not, there is no point in remembering them
 - (e.g., it helps to remember how you found a parking space near that restaurant)

Two big tasks of CBR

- Classification tasks (good for CBR)
 - Diagnosis - what type of fault is this?
 - Prediction / estimation - what happened when we saw this pattern before?
- Synthesis tasks (harder for CBR)
 - Engineering Design
 - Planning
 - Scheduling

CBR cycle



A simple example: What's a case?

- A case describes one particular diagnostic situation
- A case records several features and their specific values occurred in that situation
- A case is not a rule !!!

Technical Diagnosis of Car Faults

Figures of this example come from:

R. Bergmann, University of Kaiserslautern

CASE 1	
	<p>Problem (Symptom's)</p> <ul style="list-style-type: none">• Problem: Front light doesn't work• Car: VW Golf II, 1.6 L• Year: 1993• Battery voltage: 13,6 V• State of lights: OK• State of light switch: OK
	<p>Feature</p> <p>Value</p> <p>Solution</p> <ul style="list-style-type: none">• Diagnosis: Front light fuse defect• Repair: Replace front light fuse

		Feature	Value
C A S E 1	Problem (Symptoms)	<ul style="list-style-type: none"> • Problem: Front light doesn't work • Car: VW Golf II, 1.6 L • Year: 1993 • Battery voltage: 13,6 V • State of lights: OK • State of light switch: OK 	
	Solution	<ul style="list-style-type: none"> • Diagnosis: Front light fuse defect • Repair: Replace front light fuse 	

	Problem (Symptoms)
C A S E 2	<ul style="list-style-type: none"> • Problem: Front light doesn't work • Car: Audi A6 • Year: 1995 • Battery voltage : 12,9 V • State of lights: surface damaged • State of light switch: OK
	Solution
	<ul style="list-style-type: none"> • Diagnosis: Bulb defect • Repair: Replace front light

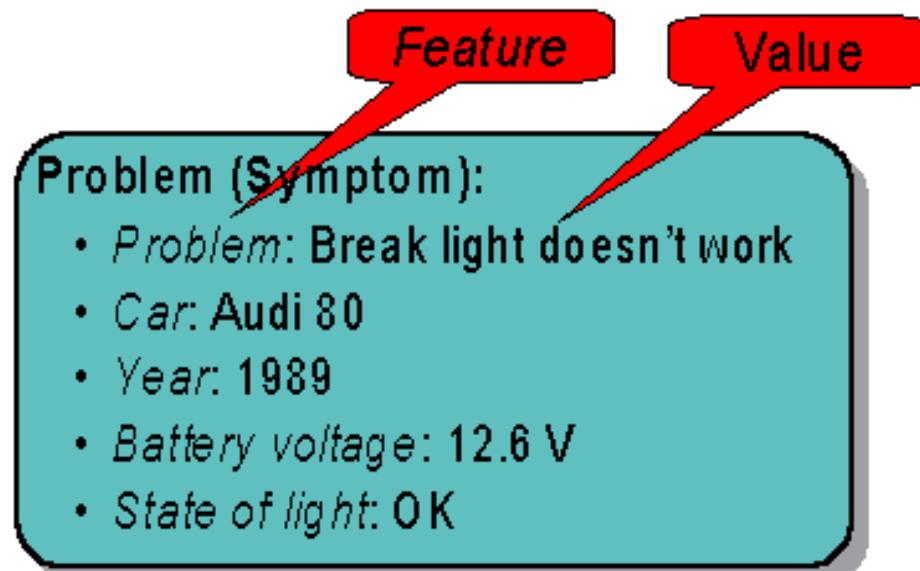
A case base with two cases:

- Each case describes a particular situation
- All cases are independent from each other

- Flat feature-value list
 - Simple case structure is sometimes sufficient for problem solving
 - Easy to store and retrieve in a CBR system
- Object Oriented representation
 - Case: collection of objects (instances of classes) in the sense of OO
 - Required for complex and structured objects
- For special tasks:
 - Graph representations: case = set of nodes and arcs
 - Plans: case = (partially) ordered set of actions
 - Predicate logic: case = set of atomic formulas
- The choice of representation is
 - Dependent on requirements of domain and task
 - Structure of already available case data

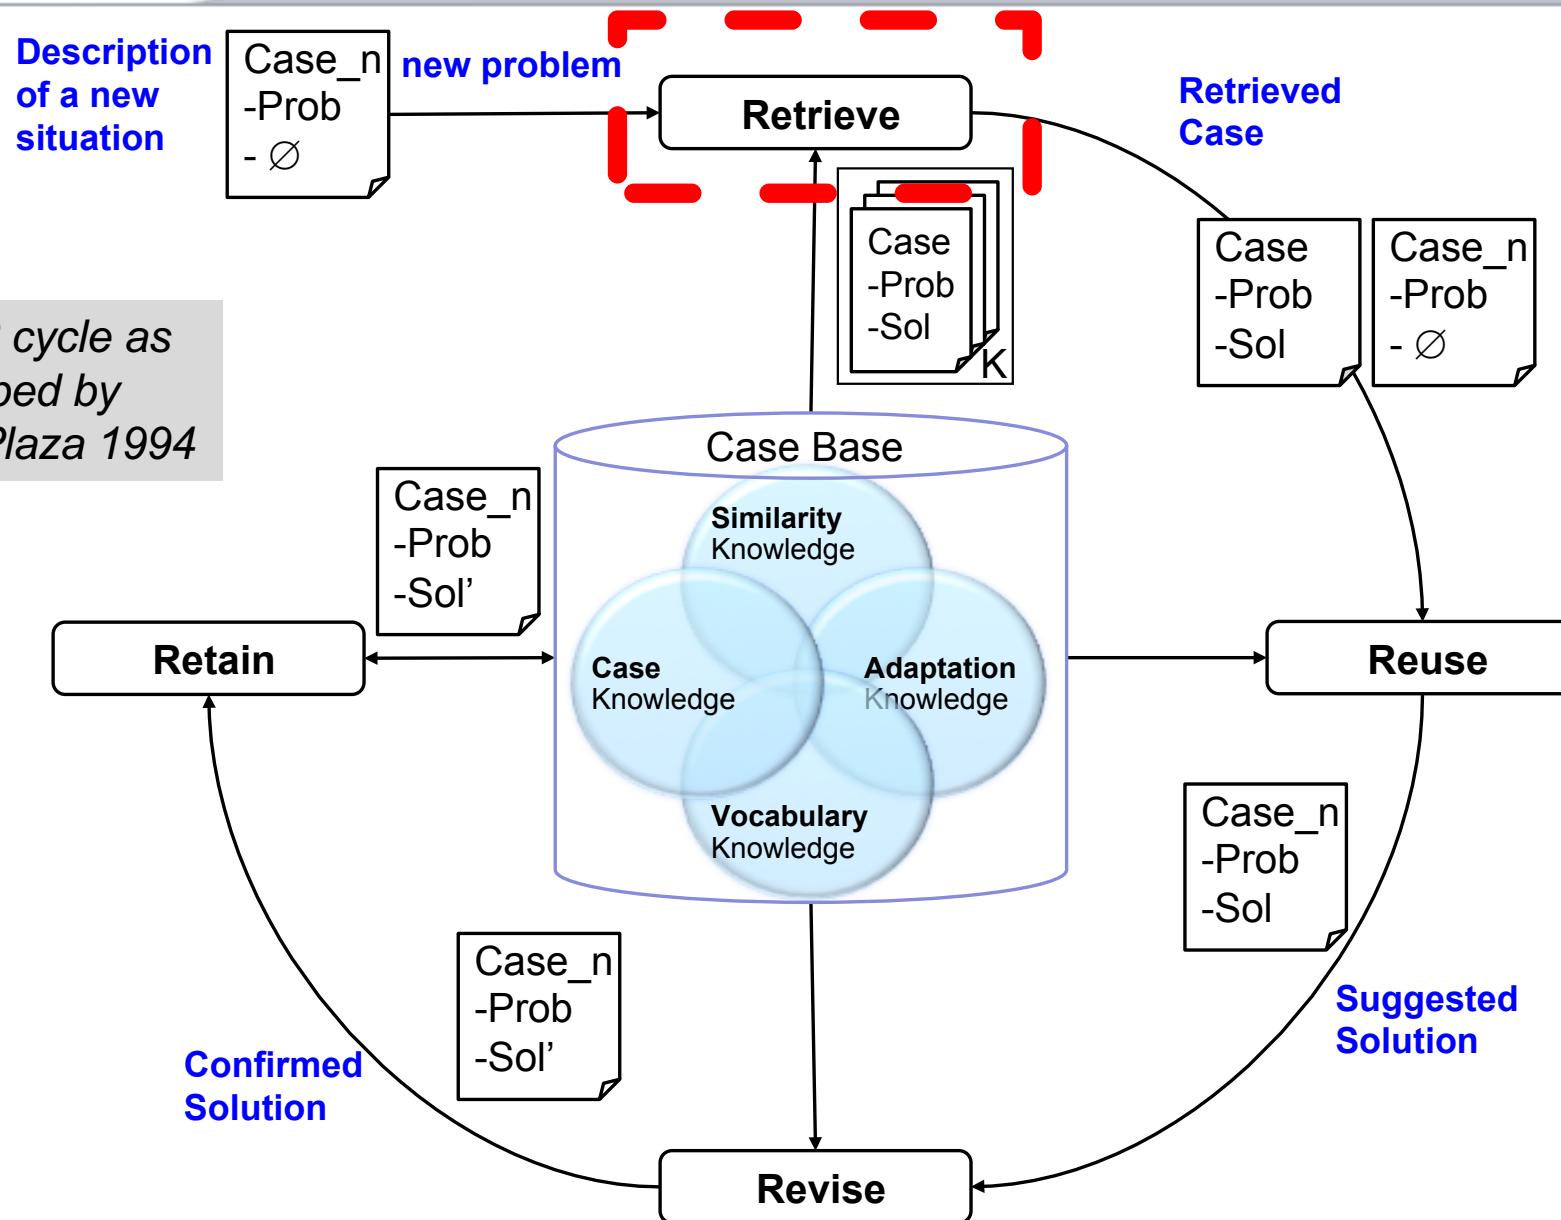
Solving a new diagnostic problem

- A new problem must be solved
- We make several observations in the current situation
- Observations define a new problem
- Not all feature values must be known
- Note: The new problem is a case without solution part



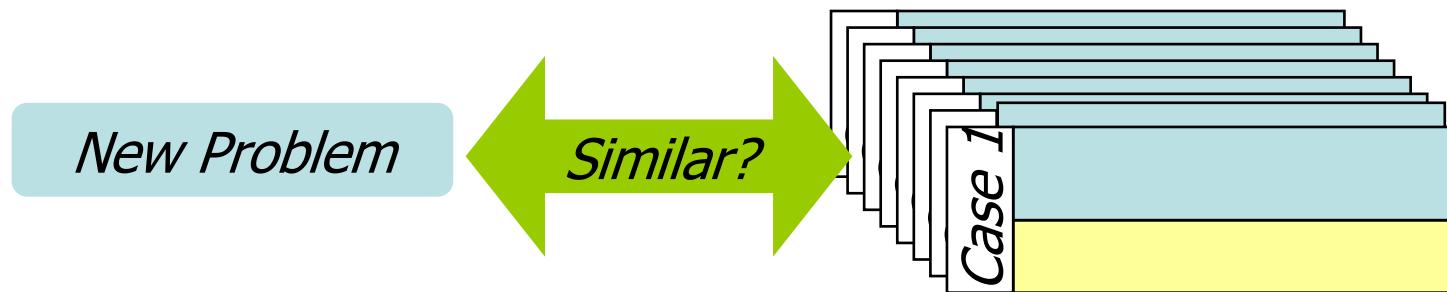
CBR cycle

The CBR cycle as described by Aamodt & Plaza 1994



Retrieving A Car Diagnosis Case

- Compare new problem to each case
- Select most similar



- Similarity is most important concept in CBR
 - When are two cases similar?
 - How are cases ranked according to similarity?
- Similarity of cases
 - Similarity for each feature
 - Depends on feature values
 - BUT: Importance of different feature values may be different

- Purpose of similarity:
 - Select cases that can be adapted easily to the current problem
 - Select cases that have (nearly) the same solution than the current problem
- Basic assumption: similar problems have similar solutions
- Degree of similarity = utility /reusability solution
- Goal of similarity modeling: provide a good approximation
 - Close to real reusability
 - Easy to compute

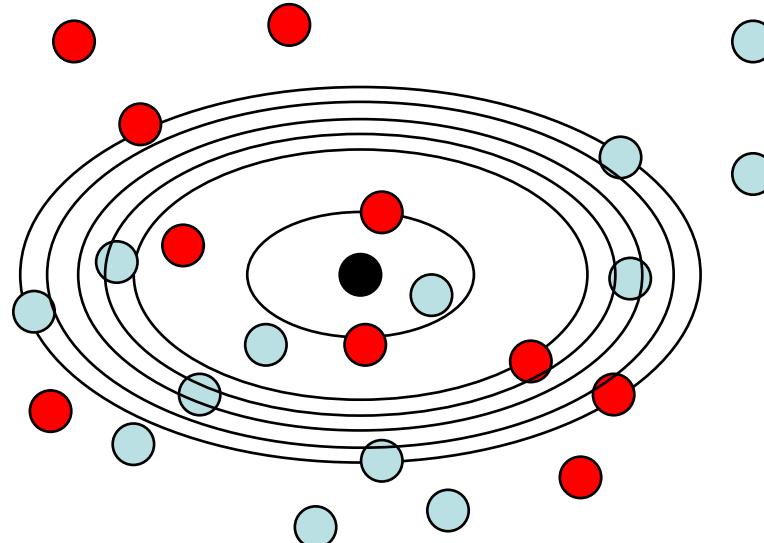
- Retrieve most similar
- k-nearest neighbour
 - k-NN
 - like scoring in bowls or curling
- Example

?

1-NN

?

5-NN



- Different approaches depending on case representation
- Similarity measures:
 - Function to compare two cases
 - sim: Case x Case → [0..1]*
 - Local similarity measure: similarity on feature level
 - Global similarity measure: similarity on case or object level
 - Combines local similarity measures
 - Takes care of different importance of attributes (weights)
 - (Sub-)Graph isomorphism for graph representations
 - Logical inferences

Similarity Computation for case

1

Problem (Symptom)

- Prob.: Break light doesn't work
- Car: Audi 80
- Year: 1989
- Battery voltage: 12.6 V
- State of lights: OK

Problem (Symptoms)

- Problem: Front light doesn't work
- Car: VW Golf II, 1.6 L
- Year: 1993
- Battery voltage: 13.6 V
- State of lights: OK
- State of light switch: OK

Solution

- Diagnosis: Front light fuse defect
- Repair: Replace front light fuse

Very important feature: weight = 6 ↪
Less important feature: weight = 1 ↪

Similarity Computation by Weighted Average

$$\text{similarity}(\text{new}, \text{case 1}) = 1/20 * [6 * 0.8 + 1 * 0.4 + 1 * 0.6 + 6 * 0.9 + 6 * 1.0] = 0.86$$

Similarity Computation for case 2

Problem (Symptom)

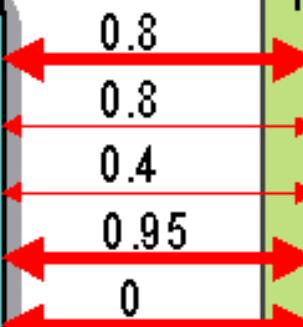
- Prob.: Break light doesn't work
- Car: Audi 80
- Year: 1989
- Battery voltage: 12.6 V
- State of lights: OK

Problem (Symptoms)

- Problem: Front light doesn't work
- Car: Audi A6
- Year: 1995
- Battery voltage : 12.9 V
- State of lights: surface damaged
- State of light switch: OK

Solution

- Diagnosis: Bulb defect
- Repair: Replace front light



Very important feature: weight = 6

Less important feature: weight = 1

Similarity Computation by Weighted Average

$$\text{similarity}(\text{new}, \text{case 2}) = 1/20 * [6*0.8 + 1*0.8 + 1*0.4 + 6*0.95 + 6*0] = 0.585$$

- Efficient case retrieval is essential for large case bases
- Different approaches depending
 - On the case representation
 - Size of the case base
- Organisation of the case base:
 - Linear lists, only for small case bases
 - Index structures for large case bases
 - Kd-trees, Retrieval nets, discrimination nets, ...
- How to store cases:
 - Databases: for large case bases or if shared with other applications
 - Main memory: for small case bases, not shared

New Problem

- Symptom: brakelight does not work
- Car: Ford Fiesta
- Year: 1997
- Battery: 9.2v
- Headlights: undamaged
- HeadlightSwitch: ?

Problem

- Case 1*
- Symptom: headlight does not work
 - ...

Solution

- Diagnosis: headlight fuse blown
- Repair: replace headlight fuse

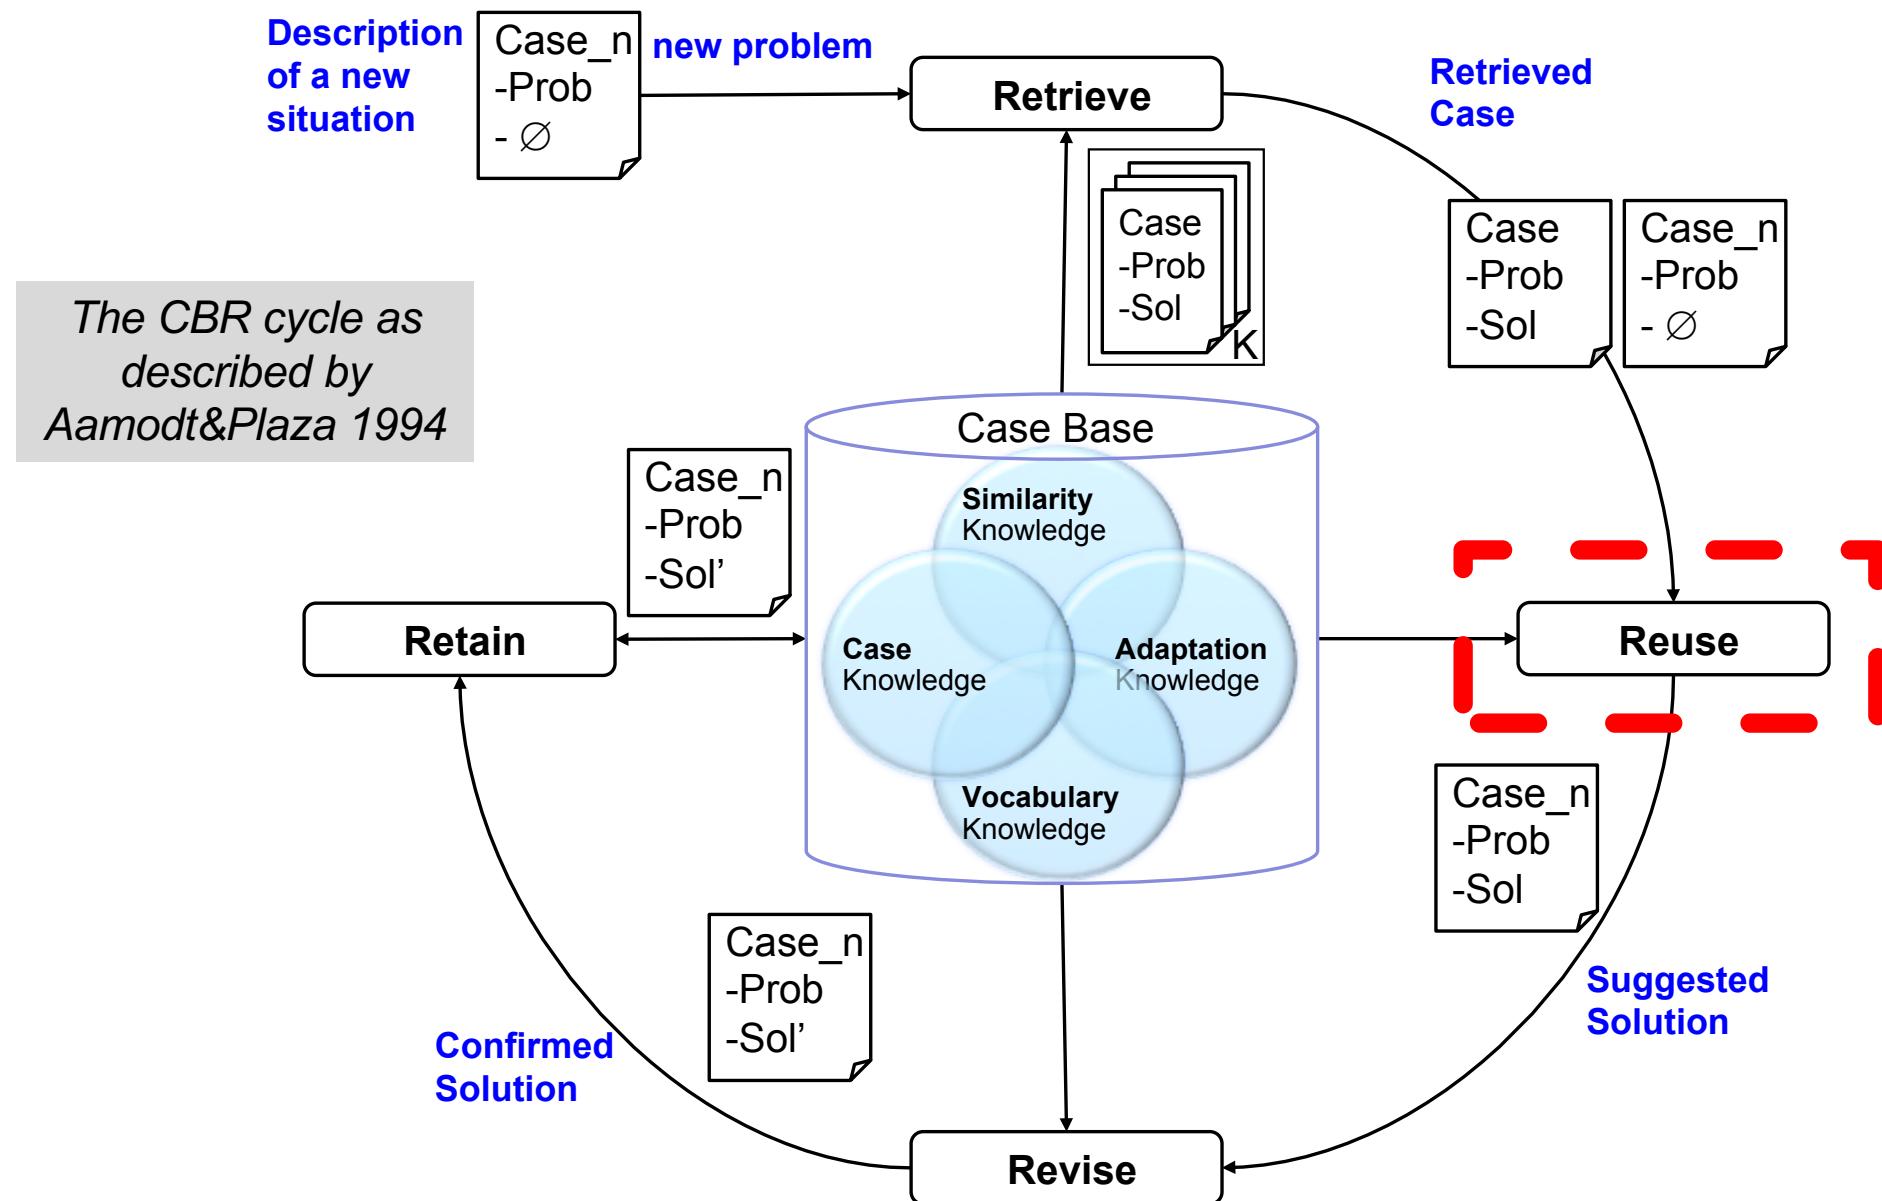
— Solution to New Problem

- Diagnosis: headlight fuse blown
- Repair: replace headlight fuse

— After Adaptation

- Diagnosis: brakelight fuse blown
- Repair: replace brakelight fuse

CBR cycle



- Single retrieved solution
 - Re-use this solution
- Multiple retrieved solutions
 - Vote/average of retrieved solutions
 - Weighted according to
 - Ranking
 - Similarity
- Iterative retrieval
 - Solve components of the solution one at a time

- Adaptation alters proposed solution:
- Null adaptation - copy retrieved solution
 - Used by CBR-Lite systems
- Manual or interactive adaptation
 - User adapts the retrieved solution (Adapting is easier than solving?)
- Automated adaptation
 - CBR system is able to adapt the retrieved solution
 - Adaptation knowledge required

- *Substitution*
 - change some part(s) of the retrieved solution
 - simplest and most common form of adaptation
- *Transformation*
 - alters the structure of the solution
- *Generative*
 - replays the method of deriving the retrieved solution on the new problem
 - most complex form of adaptation

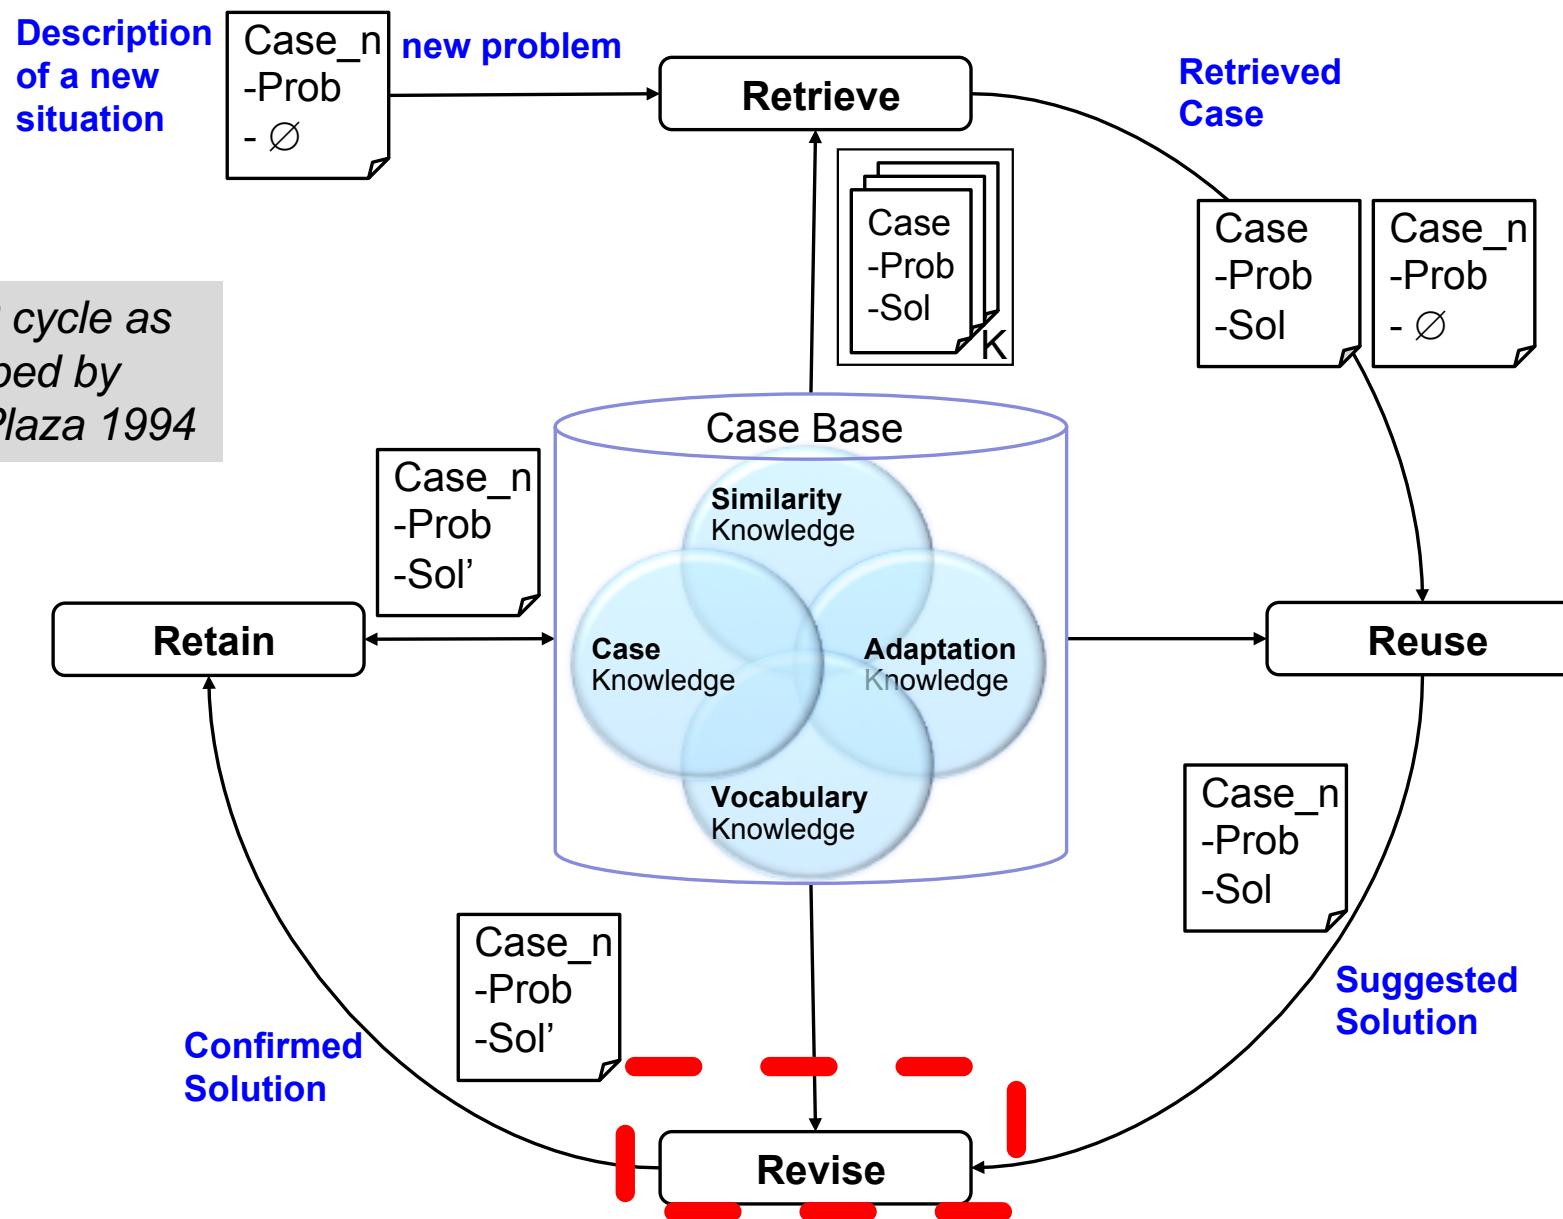
- CHEF
 - CBR system to plan Szechuan recipes
 - Hammond (1990)
- *Substitution adaptation*
 - substitute ingredients in the retrieved recipe to match the menu
 - Retrieved recipe contains beef and broccoli
 - New menu requires chicken and snowpeas
 - Replace chicken for beef, snowpeas for broccoli
- *Transformation adaptation*
 - Add, change or remove steps in the recipe
 - Skinning step added for chicken, not done for beef

- Car diagnosis example
 - Symptoms, faults and repairs for brake lights are analogous to those for headlight
 - *Substitution*: brake light fuse
- Planning example
 - Train journeys and flights are analogous
 - *Transformation*: flights need check-in step added



CBR cycle

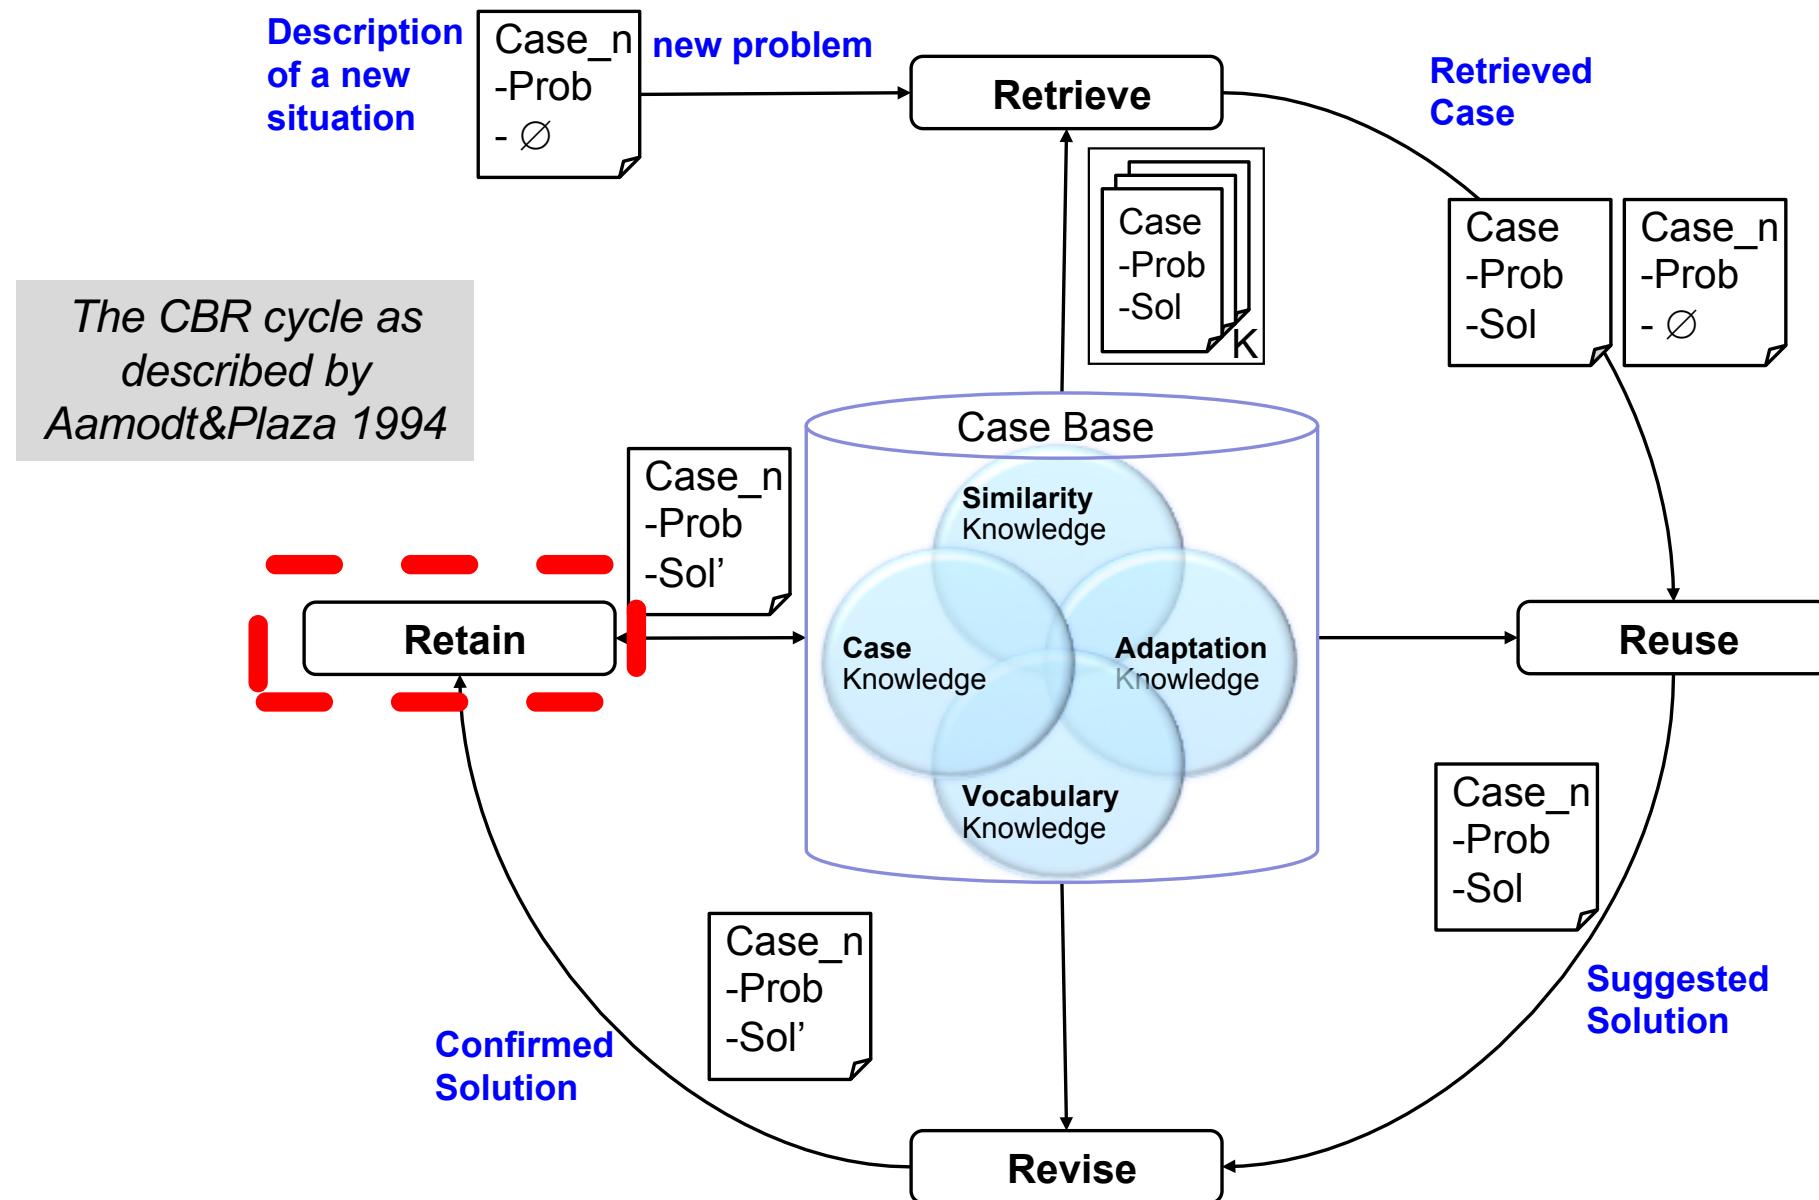
The CBR cycle as described by Aamodt&Plaza 1994



- Revise phase:
 - No revise phase
 - Verification of the solution by computer simulation
 - Verification / evaluation of the solution in the real world
- Criteria for revision
 - Correctness of the solution
 - Quality of the solution
 - Other, e.g., user preferences



CBR cycle



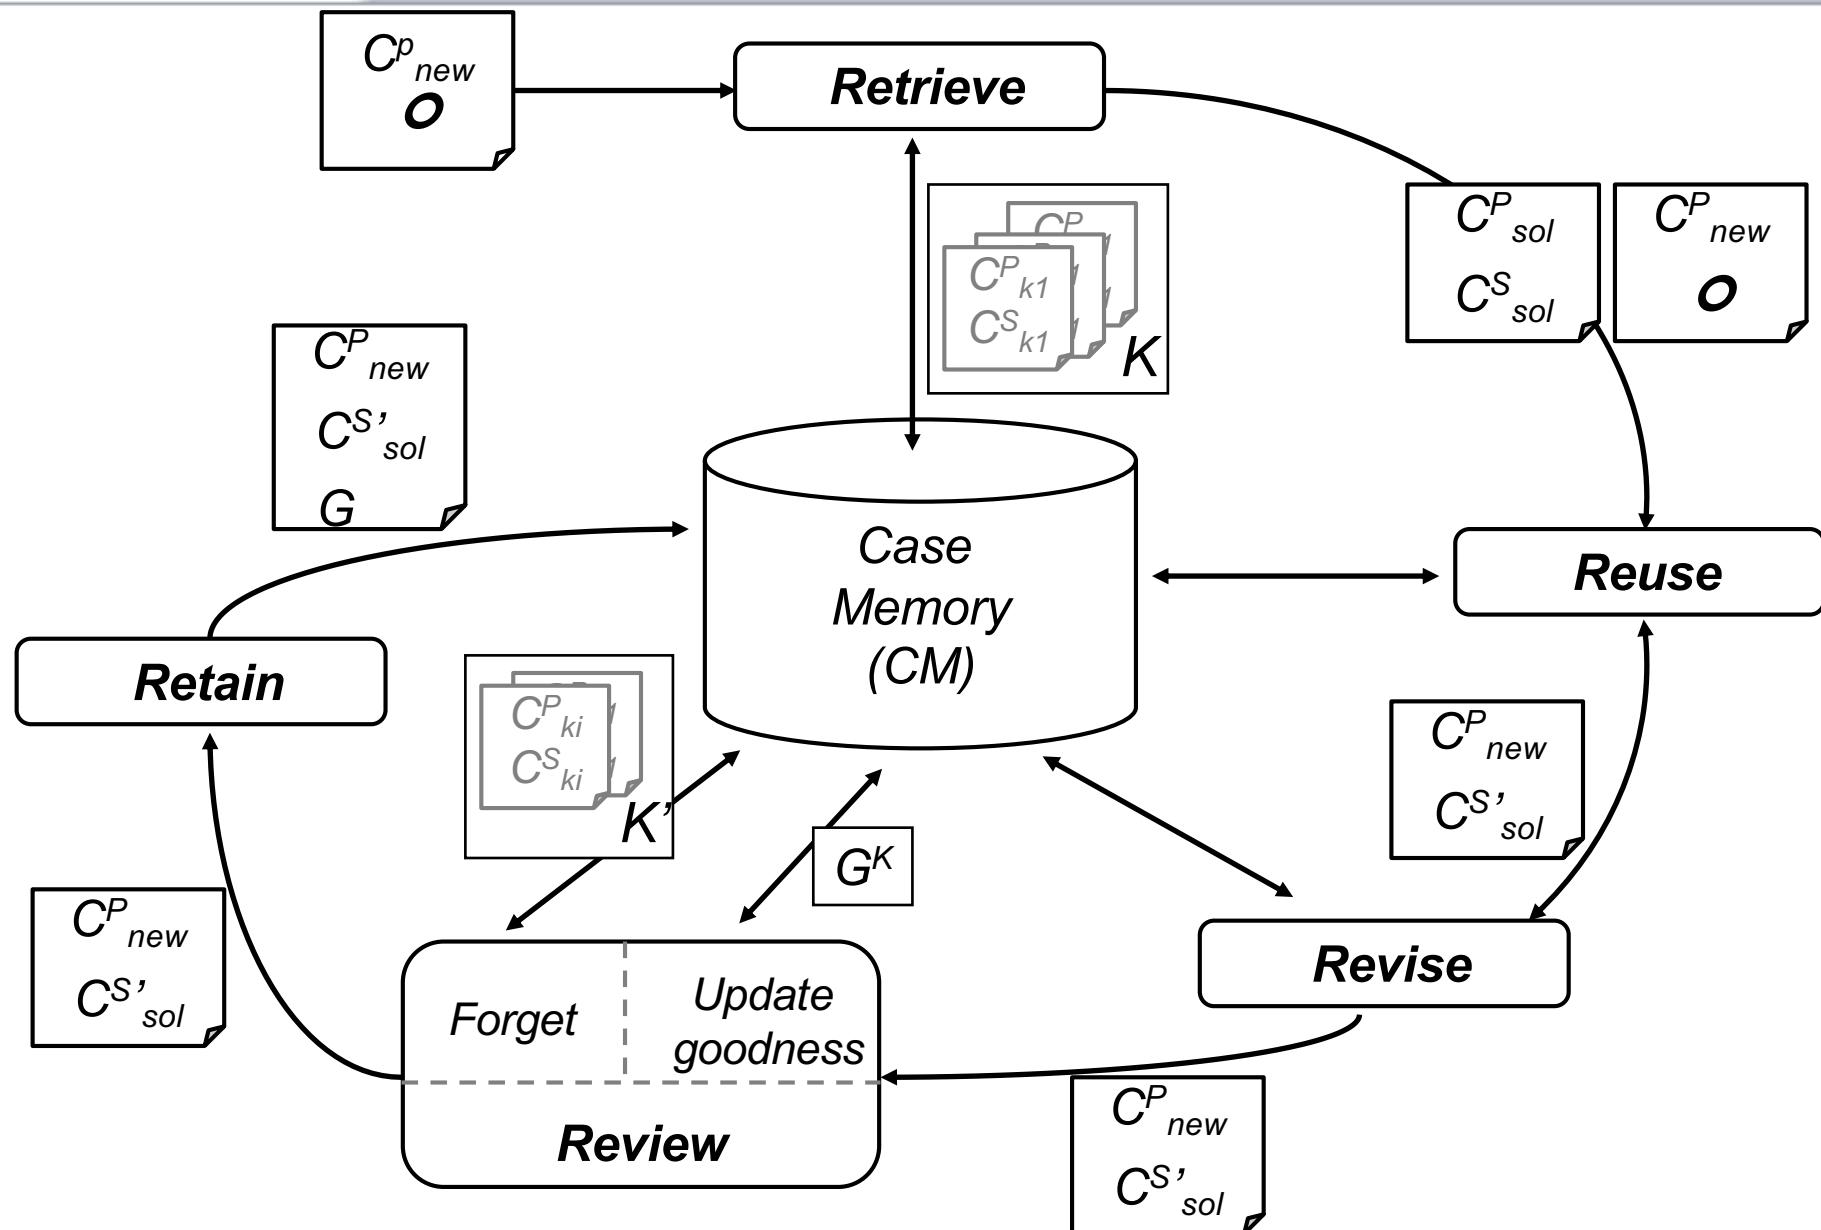
- What can be learned:
 - New experience to be retained as new case
 - Improved similarity assessment, importance of features
 - Organization / Indexing of the case base to improve efficiency
 - Knowledge for solution adaptation
 - Forgetting cases (learn to forget)
 - For efficiency or because out of date
 - Deleting an old case
 - Old is not necessarily bad
 - Does it leave a gap?



- Advantages
 - solutions are quickly proposed
 - derivation from scratch is avoided
 - domains do not need to be completely understood
 - cases useful for open-ended/ill-defined concepts
 - highlights important features
- Disadvantages
 - old cases may be poor
 - library may be biased
 - most appropriate cases may not be retrieved
 - retrieval/adaptation knowledge still needed

- Reduces the knowledge acquisition effort
- Requires less maintenance effort
- Improves problem solving performance through reuse
- Makes use of existing data, e.g., in databases
- Improve over time and adapt to changes in the environment
- High user acceptance

A CBR approach: ACBR Cycle



- CBR is a technique for solving problems based on experience
- CBR problem solving involves four phases
 - Retrieval, Reuse, Revise, and Retain
- Different techniques for:
 - Representing the knowledge, in particular, the cases
 - Realizing the four phases
- CBR has several advantages over tradition KBS
- Several applications
 - Classification, diagnosis, decision support, planning, configuration, design, ...



UNIVERSITAT DE BARCELONA



Course. Introduction to Machine Learning

Lecture 7. Lazy learning

THE END